

Note to the Grader:

Dear grader,

I have completed all the problem sets (1-5) but with the recent change in core curriculum Linear regression 3.5 solution has disappeared. I am adding the solution below:

```
import numpy as np
```

```
import pandas
```

```
from ggplot import *
```

```
"""
```

In this question, you need to:

- 1) implement the `compute_cost()` and `gradient_descent()` procedures
- 2) Select features (in the predictions procedure) and make predictions.

```
"""
```

```
def normalize_features(df):
```

```
    """
```

```
    Normalize the features in the data set.
```

```
    """
```

```
    mu = df.mean()
```

```
    sigma = df.std()
```

```
    if (sigma == 0).any():
```

```
        raise Exception("One or more features had the same value for all samples, and thus could " + \
                        "not be normalized. Please do not include features with only a single value " + \
                        "in your model.")
```

```
    df_normalized = (df - df.mean()) / df.std()
```

```
    return df_normalized, mu, sigma
```

```
def compute_cost(features, values, theta):
```

```
    """
```

Compute the cost function given a set of features / values,
and the values for our thetas.

This can be the same code as the compute_cost function in the lesson #3 exercises,
but feel free to implement your own.

```
    """
```

```
    # your code here
```

```
    m = len(values)
```

```
    sum_of_square_errors = np.square(np.dot(features, theta) - values).sum()
```

```
    cost = sum_of_square_errors / (2*m)
```

```
    return cost
```

```
def gradient_descent(features, values, theta, alpha, num_iterations):
```

```
    """
```

Perform gradient descent given a data set with an arbitrary number of features.

This can be the same gradient descent code as in the lesson #3 exercises,
but feel free to implement your own.

```
    """
```

```
    m = len(values)
```

```
    cost_history = []
```

```
    for i in range(num_iterations):
```

```
        # your code here
```

```
        Predicted_values = np.dot(features,theta)
```

```
theta = theta - (alpha/m) * np.dot((Predicted_values - values), features)
```

```
cost = compute_cost(features, values, theta)
```

```
return theta, pandas.Series(cost_history)
```

```
def predictions(dataframe):
```

```
'''
```

The NYC turnstile data is stored in a pandas dataframe called weather_turnstile.

Using the information stored in the dataframe, let's predict the ridership of the NYC subway using linear regression with gradient descent.

You can download the complete turnstile weather dataframe here:

https://www.dropbox.com/s/meyki2wl9xfa7yk/turnstile_data_master_with_weather.csv

Your prediction should have a R^2 value of 0.40 or better.

You need to experiment using various input features contained in the dataframe.

We recommend that you don't use the EXITSn_hourly feature as an input to the linear model because we cannot use it as a predictor: we cannot use exits counts as a way to predict entry counts.

Note: Due to the memory and CPU limitation of our Amazon EC2 instance, we will give you a random subet (~15%) of the data contained in turnstile_data_master_with_weather.csv. You are encouraged to experiment with this computer on your own computer, locally.

If you'd like to view a plot of your cost history, uncomment the call to plot_cost_history below. The slowdown from plotting is significant, so if you are timing out, the first thing to do is to comment out the plot command again.

If you receive a "server has encountered an error" message, that means you are hitting the 30-second limit that's placed on running your program. Try using a smaller number for num_iterations if that's the case.

If you are using your own algorithm/models, see if you can optimize your code so that it runs faster.

```
'''
```

```
# Select Features (try different features!)
```

```
features = dataframe[['rain', 'Hour','meantempi','fog','precipi']
```

```
# Add UNIT to features using dummy variables
```

```
dummy_units = pandas.get_dummies(dataframe['UNIT'], prefix='unit')
```

```
features = features.join(dummy_units)
```

```
# Values
```

```
values = dataframe['ENTRIESn_hourly']
```

```
m = len(values)
```

```
features, mu, sigma = normalize_features(features)
```

```
features['ones'] = np.ones(m) # Add a column of 1s (y intercept)
```

```
# Convert features and values to numpy arrays
```

```
features_array = np.array(features)
```

```
values_array = np.array(values)
```

```
# Set values for alpha, number of iterations.
```

```
alpha = 0.1 # please feel free to change this value
```

```
num_iterations = 100 # please feel free to change this value
```

```
# Initialize theta, perform gradient descent
```

```
theta_gradient_descent = np.zeros(len(features.columns))
theta_gradient_descent, cost_history = gradient_descent(features_array,
                                                         values_array,
                                                         theta_gradient_descent,
                                                         alpha,
                                                         num_iterations)
```

```
plot = None
```

```
# -----
```

```
# Uncomment the next line to see your cost history
```

```
# -----
```

```
#plot = plot_cost_history(alpha, cost_history)
```

```
#
```

```
# Please note, there is a possibility that plotting
```

```
# this in addition to your calculation will exceed
```

```
# the 30 second limit on the compute servers.
```

```
predictions = np.dot(features_array, theta_gradient_descent)
```

```
print theta_gradient_descent ## prints out theta/weights/coefficients
```

```
return predictions, plot
```

```
def plot_cost_history(alpha, cost_history):
```

```
    """This function is for viewing the plot of your cost history.
```

```
    You can run it by uncommenting this
```

```
        plot_cost_history(alpha, cost_history)
```

```
call in predictions.
```

If you want to run this locally, you should print the return value from this function.

```
"""
```

```
cost_df = pandas.DataFrame({
```

```
    'Cost_History': cost_history,
```

```
    'Iteration': range(len(cost_history))
```

```
})
```

```
return ggplot(cost_df, aes('Iteration', 'Cost_History')) + \
```

```
    geom_point() + ggtitle('Cost History for alpha = %.3f' % alpha )
```